


Pattern Recognition


ECE 4363 / ECE 5363

Project 2

1. Use Matlab's `quadprog()` function or Python's CVXOPT package to implement the linearly nonseparable (soft margin) SVM in its dual form and test its functionality with the data set generated as shown below. For $C = 0.1$ and $C = 100$, plot the samples, margin hyperplanes, and the decision boundary. Also, on the plot, identify and give the count of the support vectors and the misclassified samples.




First plot. $C = 0.1$. 30 Support Vectors. 8 Misclassifications. 1 Misclassification that is outside of the margins (the top red data point).



Second plot. C = 100. 30 Support Vectors. 8 Misclassifications. Thinner margin, 1 less misclassification. Visually looks like it's probably a better fit actually despite higher C usually leading towards higher overfitting.

2. Compare the computational efficiency of your implementation of SVM with one in Matlab or Python that employs the SMO approach. Present this comparison as a plot of the number of training samples versus execution time.



Plot of number of samples versus time to execute each method in seconds. I did 10 iterations at 100,200,300,...,1000 samples each. The quad prog implementation grows almost exponentially. The

other method I used is Scikit-Learn's SVM package that uses the SMO algorithm. It looks like a flat line on this scale but it's actually growing almost linearly. Much faster than my implementation (by a couple orders of magnitude).